



Sarah Guo — Pressure-Test Conversation Script

Date: Tuesday June 2, 2026 (target: afternoon call, before Wed Jun 3 operating-model send to TVP + Stillmark) **Caller:** Veronica S. Dawkins **Recipient:** Sarah Guo, Founder/Managing Partner, Conviction **Format:** 20-25 minute call. Veronica drives. Sarah answers. **Stakes:** This call decides whether the operating model gets sent to TVP and Stillmark on Wed Jun 3 as-is, or gets a 6-12 hour revision sprint Tue night before the send. **Operating reference:** `14-seed-close-pack/01-cfo-24-month-model.md` (Base + Stretch + Downside cases)

Why this call, why Sarah specifically

The 24-month operating model has a Base case (\$9.4M Y1 ARR booked, \$6.7M recognized) and a Stretch case (\$14M Y1 ARR booked, EBITDA breakeven Q1 2028). The reality-check workflow already stripped out the indefensible numbers from the original \$86M / \$65B vision. What's left in Stretch is meant to be aggressive-but-defensible. **Sarah's job in this call is to attack Stretch from an investor's vantage point and find the assumptions that don't survive 90 seconds of pressure.**

If she finds 2 or fewer, the model is ready for TVP/Stillmark Wed. If she finds 3-5, we do a focused revision Tue night before sending. If she finds 6+, we hold the send and rebuild Stretch.

Why Sarah: (a) She has done this pressure-test on 100+ pre-seed and seed models. (b) She has specific AI-platform experience post-Greylock, now at Conviction, that's directly relevant. (c) She is friend-of-the-founder enough to give an honest read, NOT polite-investor read. (d) She is NOT in the co-lead pool — so her pressure-test doesn't double-count as a diligence touch with TVP or Stillmark.

Opening (45 seconds)

"Sarah, thank you for taking this. Direct ask: I want you to attack the Stretch case of our operating model from a TVP/Stillmark vantage point and tell me what doesn't survive 90 seconds of pushback. I'm sending the full model to those two firms tomorrow morning and I'd rather find out from you tonight than from them Wednesday. I'll send you the model right after we hang up if it would help, but the questions below are designed to surface the worst objections without you having to read it first.

I'll drive. Twenty minutes. Go?"

The 12 pressure-test questions

Q1 — The Article 50 enforcement-velocity assumption

The Stretch case assumes 14% Founding Member close rate on outbound vs. 8% in Base. That delta is justified by Article 50 enforcement being VISIBLE between Aug 2 and Sep 30 — meaning at least one regulator-side enforcement action that the market reads as "this is real." If enforcement velocity tracks GDPR (18-24 months to visible action), what happens to the Stretch case?

Listen for: Does she think the Aug-Sep visible-enforcement assumption is plausible? If she pushes back hard ("never going to happen in 60 days"), Stretch needs to be derisked or pulled.

Q2 — The cohort conversion rate from Founding Member to Tier-2 (\$120K ACV)

The Base case has 30% of Founding Members converting to Tier-2 within 12 months. The Stretch case has 45%. The model has this conversion landing primarily in Q1-Q2 2027. What's a Conviction-portfolio reasonable comparable for 12-month design-partner → mid-market-expansion conversion in regulated B2B SaaS?

Listen for: She'll cite specific data from her portfolio. If 45% is way out of range, Stretch breaks. If 30% Base is generous, Base breaks too.

Q3 — The Big-4 advisory recommendation timing

Stretch assumes the first Big-4 co-authored advisory citing "LP-Conformant" lands by W19 (Oct 6, 2026). That's ~4 months from the W7 (Jul 21) Big-4 working group invitation. Have you seen Big-4 advisory recommendations move in 4 months in your portfolio, or does that usually take 12+ months?

Listen for: If she says "12-18 months minimum, never seen 4 months," we delete the Big-4 recommendation line from the H1 2027 revenue stack in Stretch.

Q4 — The IETF SCITT WG adoption timing

Stretch assumes IETF SCITT WG adoption confirmed by W23 (Nov 3 — IETF 121 Dublin) for `draft-dawkins-scitt-ai-article50-01`. The fallback is CEN-CENELEC JTC 21 NWIP via DIN as alternate standards path. From your view of the IETF / standards-body world, is "adopted within 6 months of first IETF posting" achievable for a working-group-not-individual-submission draft on a sensitive AI-policy topic?

Listen for: She may not be an IETF person directly but Conviction's portfolio touches enough of this. If she signals "6 months is fast" we don't change the model; if she signals "you'll be lucky to get adoption in 18 months," we add a derisking line to Stretch.

Q5 — The Strategic Beta Partner \$1M pricing — Riot specifically

The Stretch case includes Riot Games Strategic Beta Partner at \$1M (with a \$1M discount on a \$2M list, accounted as \$1M cash per ASC 606). Tencent ownership of Riot creates EU-press risk we're mitigating with counterweight anchors. Two questions: (a) Is \$1M Strategic Beta Partner pricing defensible for a pre-launch protocol with no public production deployments? (b) Does the Tencent risk get raised by TVP or Stillmark independently, and how do we frame Riot in the model summary without triggering it?

Listen for: (a) Pricing — if she says "way underpriced" we leave it; if "way overpriced," we flag in the model. (b) Tencent — she'll have a sophisticated take on whether to include Riot in the customer-roster cohort visible to investors at all, or whether to defer disclosure to post-close.

Q6 — The Foundation/Inc. expense allocation

The model splits expenses 60/40 Inc./Foundation in Q3 2026, shifting to 75/25 by Q1 2027 as commercial scales. Will Series A diligence (Qatalyst-led) accept this split or demand restatement? What's the bar?

Listen for: She has seen Foundation/Inc. dual-entity structures (Linux Foundation precedent, OpenAI early-stage). If she says "diligence will want every Foundation-related expense pre-audited by Adler & Colvin," we update the Series A diligence pack timeline.

Q7 — The seed tranche-trigger clause

\$10M @ \$45M post + \$5M tranche-conditional @ \$60M post Sep 30 on three milestones (2 FSI design partners + audit memo + IETF or CEN advance). Have you seen tranche-trigger SAFEs of this shape close in 25 days with four co-leads? What's the failure mode?

Listen for: She'll either say "totally doable" or "the three-milestone clause is what will hold up the close." If the latter, we may need to renegotiate milestone count down to 2.

Q8 — The Founding Member outbound capacity

The model assumes Veronica drives outbound through Aug 15 (when EU Enterprise AE hire onboards). Stretch has 25 signed Founding Members by Aug 31; Base has 15. Is 15-25 deals founder-driven in 8 weeks defensible for a pre-launch protocol? If not, the cash recognition curve shifts right by 30-60 days and Stretch's Q4 2026 cash position degrades.

Listen for: She knows founder-led sales velocity better than anyone. If she says "10 deals is the absolute ceiling," Stretch breaks.

Q9 — The DACH press credibility assumption

Stretch assumes Handelsblatt + FAZ coverage drives 30% of German FSI inbound in Sep-Dec 2026. Brunswick (or eq.) gets us coverage. But: does Handelsblatt actually cover EU AI compliance startups at that velocity? Or do they cover incumbents (Allianz, DZ Bank, Deutsche Telekom) on the same topic and skip startups?

Listen for: If she says "Handelsblatt covers the incumbents, not the startup," the German-FSI inbound assumption needs to be pulled or routed differently.

Q10 — The Series A teaser deck timing

Operating model assumes Series A teaser deck ready by Nov 29 (W26), Qatalyst engagement signed same week, term sheets by W27. From your view of Conviction's deal cycle, is Series A signed in 6 weeks from teaser-deck-out for a \$30M / \$180M post round? Or is the realistic timeline 12-16 weeks?

Listen for: If she says 12-16 weeks, the Y2 cash trajectory needs to assume Series A doesn't land until Q2 2027, which means Hercules draw becomes mandatory not optional.

Q11 — The Big Single Worst Number

If you had to point at ONE number in the Stretch case and say "that's the one TVP or Stillmark will probably push back on hardest," which number would it be? And what's the smallest change that would make it defensible?

Listen for: This is the meta-question. Her answer tells us the single highest-priority pre-Wed-send revision.

Q12 — The "Would You Lead?" Question

If you were leading the seed round at \$45M post with the same tranche structure I described, what would your covenants be? What would you NOT cover or carve out from standard YC SAFE language?

Listen for: She probably won't lead (Conviction generally leads Series A, not seed). But her answer reveals what TVP/Stillmark will probably ask for. Specifically: pro-rata rights size, information rights cadence, board observer seat, MFN clauses, Founder Friendly Standard SAFE compliance vs. modifications.

Closing (60 seconds)

"Sarah, thank you. Two specific things I need from you before I hang up:

1. If you had to score the model 1-10 on 'investor-ready for Wed Jun 3 send to TVP and Stillmark,' what's the number?
2. If the number is below 8, what's the SINGLE change I should make tonight before sending?

And one ask: if Conviction or anyone in your network would be a useful warm intro for Series A six months from now, I'd value being on the calendar in October. No ask before that — I just want the relationship to be real when the diligence window opens."

What I do AFTER the call (Tue night, 60-90 min sprint)

Sarah's signal	My action Tue night
Score 8-10, ≤2 objections	Send model to TVP + Stillmark Wed 09:00 PDT as-is
Score 6-7, 3-5 objections	Tue night focused revision sprint on the 3-5 most critical lines; send Wed 11:00 PDT after re-review
Score ≤5, 6+ objections	HOLD send. Revise overnight + Wed AM. Send Thu 09:00 PDT with a re-pressure-test from Mark Beeston or Lars Holmquist before send

Hard constraints on the call

- DO NOT pitch. This is a pressure-test, not a fundraising touch.
- DO NOT ask Sarah for a check. She's not in the co-lead pool by design.
- DO NOT name specific co-lead positioning details (e.g., "Stillmark said X") — her relationship to the seed conversation must stay outside-the-room.
- DO NOT reveal the FAIL-branch Brunswick prep package (KS3). The Tencent risk on Riot is a strategic concern, not an open-call topic.
- DO NOT reveal the eight C-corrections — they're operational hygiene, not investor-facing material.
- DO mention: the IETF draft confirmation on Datatracker, the Foundation Inc. dual-entity structure, Audit engagement letters returned by Fri Jun 5, DIN application targeted Jun 15.

What to send Sarah BEFORE the call (if she agrees to take it Tue afternoon)

Send 1 hour before call: - 14-seed-close-pack/01-cfo-24-month-model.md (full) - One paragraph framing: "Pressure-test the Stretch case. Find the 2-5 numbers TVP or Stillmark will push back on. Twenty minutes."

DO NOT send: - The reality-check artifact (12-premortem/05-EXPLOSIVE-ADOPTION-REALITY-CHECK.md) — that's our internal reckoning, not for external eyes yet - The Atomic Explosion master plan — too much information for a 20-min call - The counterweight anchor target list — not Sarah's lane

If Sarah is unavailable Tue PM

Backup pressure-testers in priority order: 1. Mark Beeston (Fund Recs founder, EU-FSI-deep) — alternate; he sees the EU side better than Sarah does 2. Lars Holmquist (Stack AV CEO, ex-Volvo Trucks; not from Conviction, but operationally-deep) 3. Bill Janeway (Warburg Pincus, author of *Doing Capitalism in the Innovation Economy*) — slowest to respond, highest signal if he engages

If all three are unavailable, the model goes to TVP + Stillmark Wed 09:00 PDT as-is. The cost of a 24-hour pressure-test delay is higher than the cost of a slight model over-presentation.

Document boundary: This document is internal strategic planning material for LedgerProof Foundation (501(c)(3) in formation, Delaware) and LedgerProof Inc. (Delaware C-corp). It contains assumptions, forecasts, and decisions that may change as conditions evolve. Numbers labeled BASE / STRETCH / DOWNSIDE are scenarios, not forecasts. Distribution is limited to Foundation board, Inc. officers, and named diligence counterparties under NDA or equivalent.

Confidentiality: Do not redistribute without consent of the founder. Bitcoin-anchored governance receipts apply only to ratified Foundation actions; this document is a planning artifact and is NOT receipt-anchored.